

# Parallel Computing for Science & Engineering SSC 374/394c

Spring 2015

Instructors:

Victor Eijkhout, Research Scientist, TACC

Cyrus Proctor, Research Associate, TACC



THE UNIVERSITY OF TEXAS AT AUSTIN  
**Texas Advanced Computing Center**

---

# The ideas of parallel programming

- As illustrated by Conway's *Game of Life*
- <http://youtu.be/C2vgICfQawE>
- This has the same structure as certain important applications (e.g., PDEs) but requires no math to explain.
- Note: this is about parallel *programming*, not so much about parallel *hardware*

# How do you code this?

- First the function for updating a single cell

```
def life_evaluation( cells ):
    # cells is a 3x3 array
    count = 0
    for i in [0,1,2]:
        for j in [0,1,2]:
            if i!=1 and j!=1:
                count += cells[i,j]
    cells[1,1] = life_count_evaluation(count )

def life_count_evaluation( cell,count ):
    # big if statement
```

# How do you code this?

- Now to update the whole board
- One time-stepping loop
- Two loops for the board

```
life_board.create(final_time,N,N)

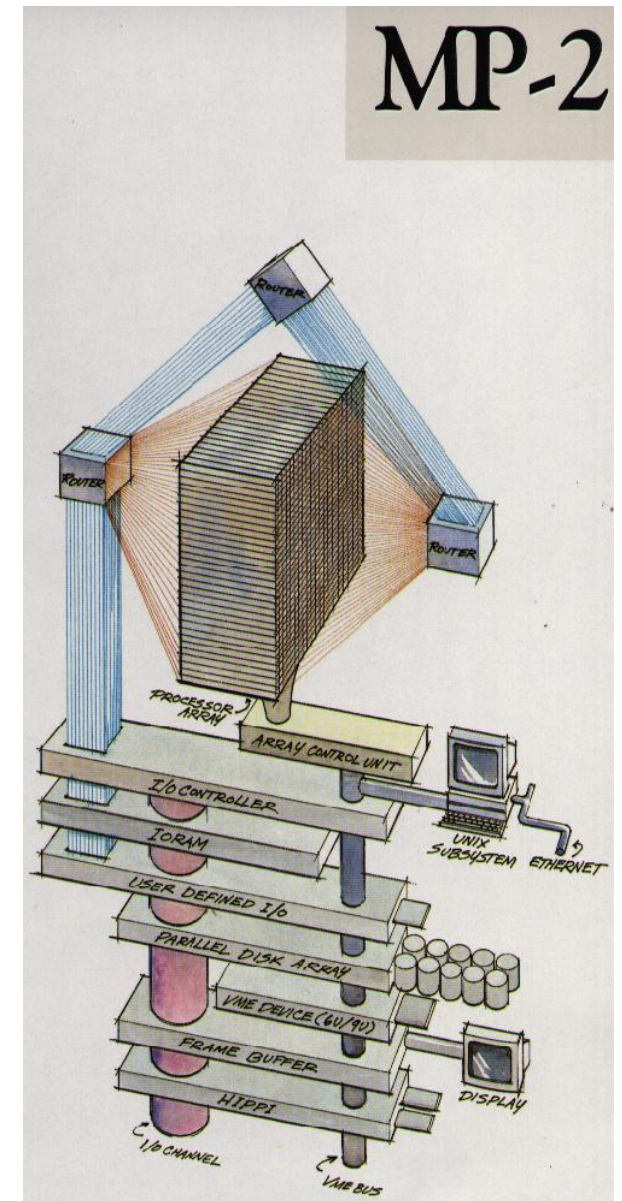
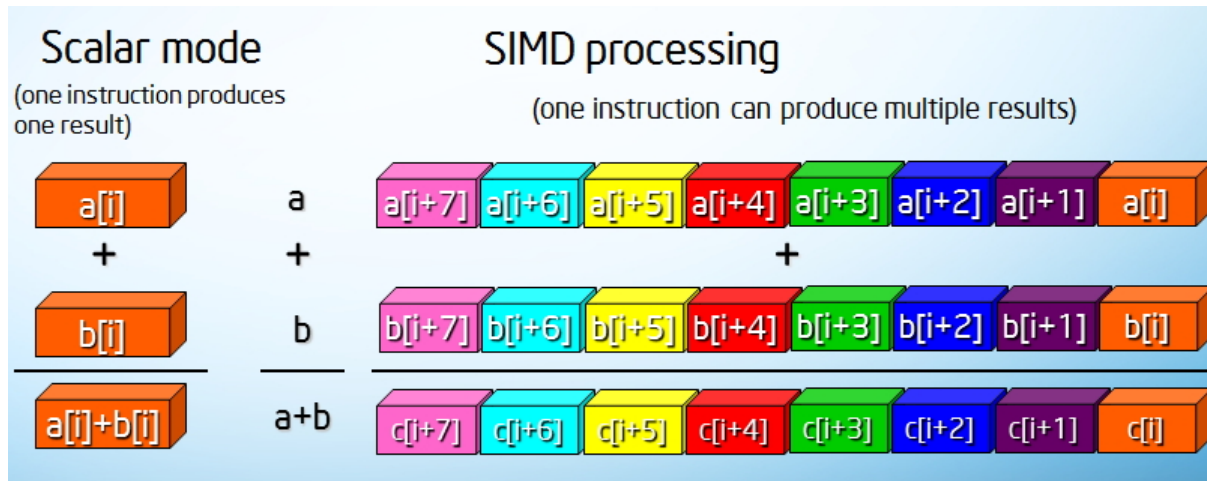
for t in [0:final_time-1]:
    for i in [0:N-1]:
        for j in [0:N-1]:
            life_board[t+1,i,j] :=
                life_evaluation
                    (life_board[t,i-1:+1,j-1:+1] )
```

# Where does parallelism come from?

- The program text specifies a sequence of operations
- However, some operations can be done in any order
- => so they can also be done simultaneously
- There are no compilers that recognize this, so you have to code it by hand and that's what you will learn here....

# Data parallelism

- Independent data items, each undergoing the same operation
- Then: “array processors”
- Now: vector instructions



# SIMD parallelism

- Single Instruction Multiple Data
- Simplified instruction handling: only one instruction fetch/decode/whatever for multiple data items
- Need to have many independent operations (examples?)
- Data storage may need to be regular

# Example: GPUs

- Graphical Processing Units are SIMD-like (not completely lockstep)
- Programmed in CUDA:  
kernel contains sequential code,  
kernel is executed in parallel

```
kerneldef life_step( board ):  
    i = my_i_number()  
    j = my_j_number()  
    board[i,j] = life_evaluation  
                    ( board[i-1:i+1,j-1:j+1] )  
  
for t in [0:final_time]:  
    <<N,N>>life_step( board )
```



# Parallel programming may mess up your code!

- Parallelism on the instruction level:  
innermost loop
- Sometimes loop exchange needed

```
for i=1,N:  
  for j=1,N:  
    count = 0  
    for h in [-1,0,1]:  
      for v in [-1,0,1]:  
        count += cell[i+v,j+h]
```

```
for i=1,N:  
  for j=1,N:  
    count[i,j] = 0  
  for h in [-1,0,1]:  
    for v in [-1,0,1]:  
      for i=1,N:  
        for j=1,N:  
          count += cell[i+v,j+h]
```

# Minimal intervention: loop parallelism

- Loops are an important source of parallelism
- Parallelize by indicating what loops parallel

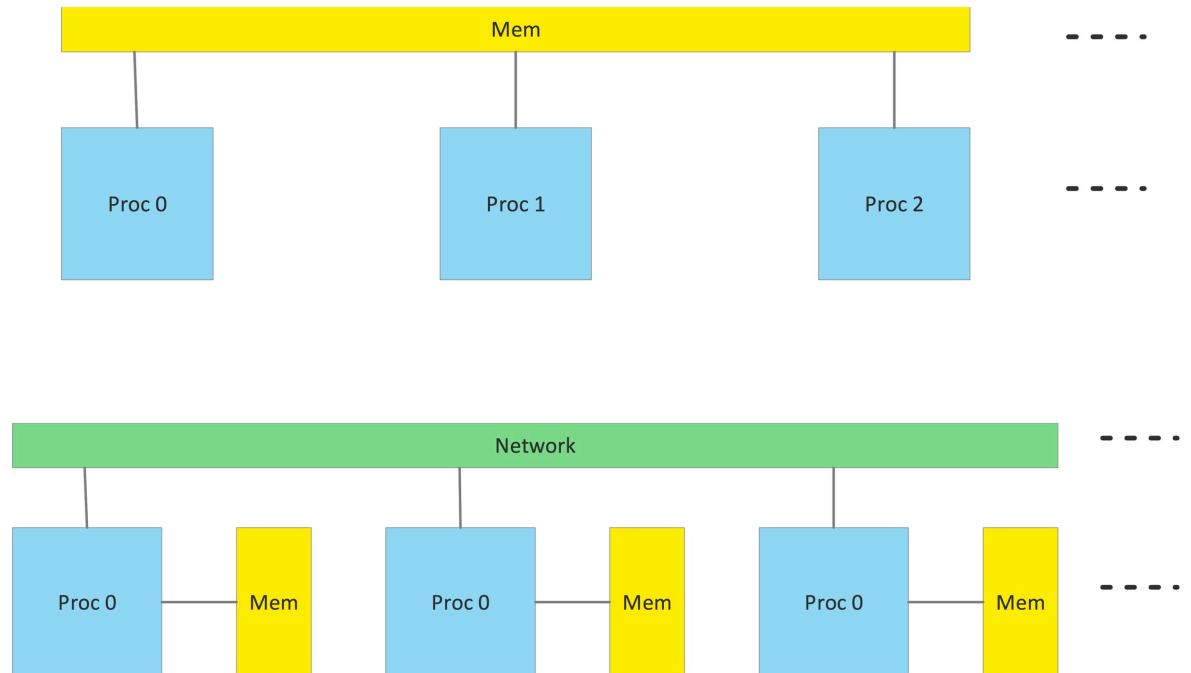
```
def life_generation( board,tmp ):  
    # OMP parallel for  
    for i in [0:N-1]:  
        for j in [0:N-1]:  
            tmp[i,j] = board[i,j]  
    # OMP parallel for  
    for i in [0:N-1]:  
        for j in [0:N-1]:  
            board[i,j] = life_evaluation  
                ( tmp[i-1:i+1,j-1:j+1] )
```

# Granularity of parallelism

- So far: independence of single operations / single data points:  
*fine-grained parallelism*
- Locality: points close together should be handled by the same processor
- Process the board by lines or subparts:  
*coarse-grained parallelism*

# Why coarse-grained parallelism?

- Shared memory: every processor can find every data item
- Distributed memory: some data is local, other not
- Locality




# What does distributed memory look like?

- Stampede
- 160 cabinets,  
6400 nodes,  
500k cores....



# How do you program distributed memory?

- Explicit message passing



```
p = my_processor_number()

high_line = MPI_Receive(from=p-1,cells=N)
low_line = MPI_Receive(from=p+1,cells=N)

tmp_line = my_line.copy()
my_line = life_line_update(high_line,tmp_line,low_line,N)
```

# No, really.....

- You can't receive without someone else sending
- But the someone else is running the same program...
- Single Program Multiple Data:  
each processor runs the same program,  
just on different data:
- Execution differs in:  
loop bounds,  
branches of conditionals

# Two-sided message passing

- Everyone does both send and receive calls
- Attempt at coding this:

- And even this is not correct

```
p = my_processor_number()

# send my data
my_line.MPI_Send(to=p-1,cells=N)
my_line.MPI_Send(to=p+1,cells=N)

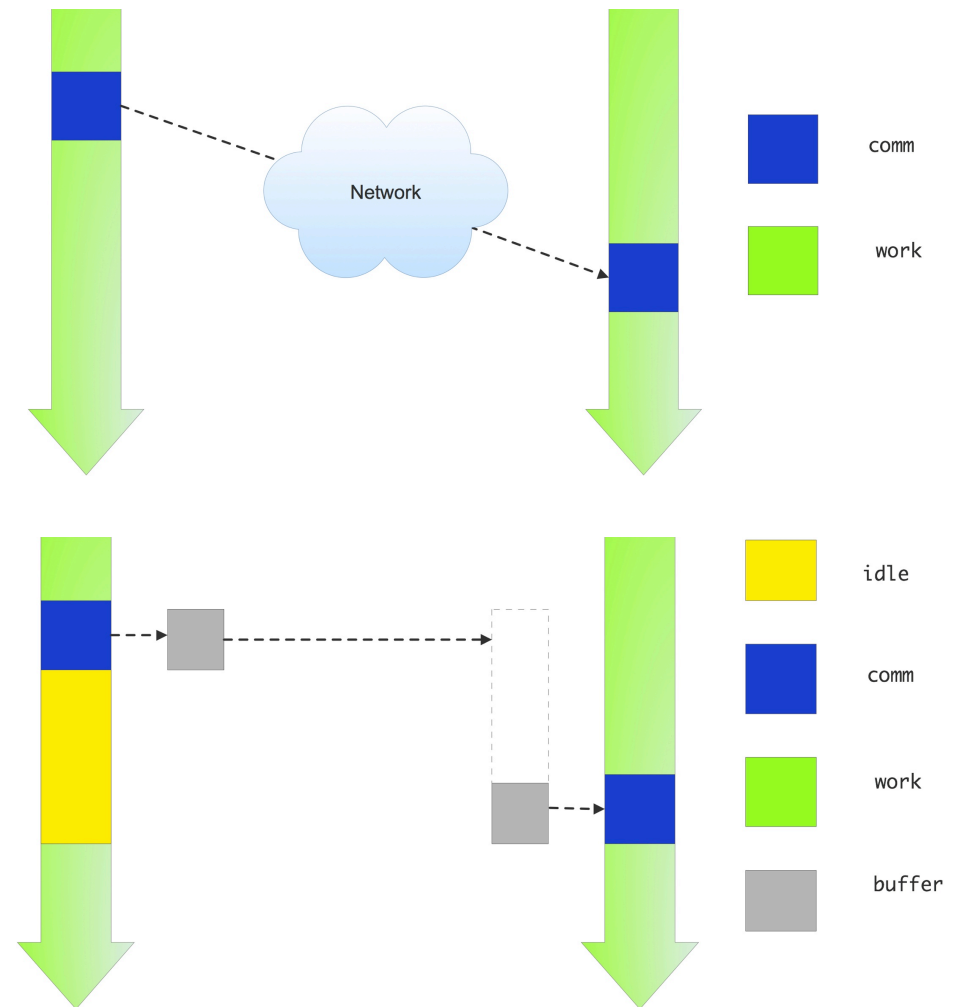
# get data from neighbours
high_line = MPI_Receive(from=p-1,cells=N)
low_line = MPI_Receive(from=p+1,cells=N)
tmp_line = my_line.copy()

# do the local computation
my_line = life_line_update
                (high_line,tmp_line,low_line,N)
```



# Blocking communication

- Data has to be somewhere
- You can only send if someone else receives
- Deadlock possible if everyone is sending, no one is receiving




# Task parallelism

- Think about instructions rather than data
- In the Game of Life there are  $N^2T$  updates
- How independent are they?

# Life updates dependencies

- Cell needs a box around it
- Each cell in the box needs....
- => Cone of influence



# Task scheduling

- User indicates dependencies
- Algorithm under the hood matches available tasks to available processors/cores

```
while there_are_tasks_left():  
    for r in running_tasks:  
        if r.finished():  
            for t in scheduled_tasks:  
                t.mark_input_available(r)  
            t = find_available_task()  
            p = find_available_processor()  
            schedule(t,p)
```

```
for t in [0:T]:  
    for i in [0:N]:  
        for j in [0:N]:  
            task( id=[t+1,i,j],  
                prereqs=[  
                    [t,i,j],  
                    [t,i-1,j],  
                    [t,i+1,j]  
                    # et cetera  
                ] )
```

# Summary

- SIMD / vector parallelism:  
very fine-grained  
vector instructions
- Loop-based parallelism:  
OpenMP directives
- Tasks:  
OpenMP tasks, medium grain
- Message passing:  
MPI, coarse grain