

PCSE Lecture 10

MPI Data Structures and Communicators

Cyrus Proctor
Victor Eijkhout

April 7, 2015

MPI Datatypes

- MPI offers a suitable collection of basic datatypes to use in messages
- Data to be communicated must be in contiguous memory locations
- Data must consist of just one basic type and no other

Solution? **MPI Derived Datatypes!**


- Derived datatypes can be built up recursively
- They can be created conditionally, at runtime
- The associated packing and unpacking are done for you automatically

Motivational Example

- Create 2D matrix (`matrix[NROWS][NCOLS]`)
(`matrix(NROWS)(NCOLS)`)
- First, send entire 2D matrix from process a to process b
- Second, send one row of 2D matrix from process a to process b
- Third, send one column of 2D matrix from process a to process b


Fortran Language

```
real*8 sa, sb
real*8 sc, d1(5), d2(5)
real*8 a(3,3)
```



C Language

```
double sa, sb;
double sc, d1[5], d2[5];
double a[3][3];
```



C – Send Full 2D Matrix

```
// Sending a full matrix from process 0 to process 1
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };

MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
MPI_Comm_rank(MPI_COMM_WORLD,&rank);

// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Send/Recv
tag1 = 1; dest=1; src=0;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i...\n\n",src,dest);
}
if(rank == 0){
    MPI_Send(&matrix[0][0],NROWS*NCOLS,MPI_DOUBLE,
            dest,tag1,MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&matrix[0][0],NROWS*NCOLS,MPI_DOUBLE,
            src,tag1,MPI_COMM_WORLD,&Stat);
}

// Print matrix after Send/Recv
...
```

Output – 2 Processes

```
(Rank 0): Matrix:
0  1  2  3
4  5  6  7
8  9 10 11
(Rank 1): Matrix:
0  0  0  0
0  0  0  0
0  0  0  0
```

Completing Send/Recv from 0 to 1...

```
(Rank 0): Matrix:
0  1  2  3
4  5  6  7
8  9 10 11
(Rank 1): Matrix:
0  1  2  3
4  5  6  7
8  9 10 11
```

C – Send One Row (default)

```
// Default way of sending contiguous one row of 2D array
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };

MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
MPI_Comm_rank(MPI_COMM_WORLD,&rank);

// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Send/Recv
tag1 = 1; dest=1; src=0; row=1;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of row %i...\n\n",src,dest,row);
}
if(rank == 0){
    MPI_Send(&matrix[row][0],NCOLS,MPI_DOUBLE,
            dest,tag1,MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&matrix[row][0],NCOLS,MPI_DOUBLE,
            src,tag1,MPI_COMM_WORLD,&Stat);
}

// Print matrix after Send/Recv
...
```

Output – 2 Processes

```
(Rank 0): Matrix:
0  1  2  3
4  5  6  7
8  9 10 11
(Rank 1): Matrix:
0  0  0  0
0  0  0  0
0  0  0  0
Completing Send/Recv from 0 to 1 of
row 1...

0
(Rank 0): Matrix:
0  1  2  3
4  5  6  7
8  9 10 11
(Rank 1): Matrix:
0  0  0  0
4  5  6  7
0  0  0  0
```

Derived types

Three main classifications

- **Contiguous Arrays** (easy to use)
 - send contiguous blocks of the same datatype
- **Noncontiguous Vectors** (relatively easy to use)
 - send noncontiguous blocks of the same datatype
- **Abstract types** (more difficult)
 - send C or Fortran 90 structures

Derived types

- **Elementary:** MPI names for language types
- **Contiguous:** Array with stride of one
- **Vector:** Array separated by constant stride
- **Hvector:** Vector, with stride in bytes
- **Indexed:** Array of indices (like gatherv)
- **Hindexed:** Indexed, with displacements in bytes
- **Struct:** General mixed types (C structs etc.)
- Pack and Unpack

Derived types, how to use them

- Three step process
- Define the type (e.g.)
 - MPI_Type_contiguous** for contiguous arrays
 - MPI_Type_vector** for noncontiguous arrays
 - MPI_Type_struct** for structures
- Commit the type
 - Tells MPI when to compile an internal representation
 - MPI_Type_commit (... my_type...)**
- Use in normal communication calls
 - MPI_Send(data, count, my_type, dest, tag, comm ...)**
- Free space when done:
 - MPI_Type_free**

Contiguous Type Example #1

C – Send One Row (MPI_Type_contiguous)

```
// Using MPI_Type_contiguous to send one row of 2D array
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };
// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Create derived contiguous type
MPI_Datatype MY_MPI_ROW;
MPI_Type_contiguous(NCOLS, MPI_DOUBLE, &MY_MPI_ROW);
MPI_Type_commit(&MY_MPI_ROW);
// Send/Recv
tag1 = 1; dest=1; src=0; row=1;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of row %i...\n\n", src, dest, row);
}
if(rank == 0){
    MPI_Send(&matrix[row][0], 1, MY_MPI_ROW,
             dest, tag1, MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&matrix[row][0], 1, MY_MPI_ROW,
             src, tag1, MPI_COMM_WORLD, &Stat);
}

// Print matrix after Send/Recv
...
```

Output – 2 Processes

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	0	0	0
0	0	0	0
0	0	0	0

Completing Send/Recv from 0 to 1 of
row 1...

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	0	0	0
4	5	6	7
0	0	0	0

Contiguous Type Example #2

C – Send One Row from 2D Array to 1D Array

```
// Using MPI_Type_contiguous to send one row of 2D array
// to 1D array
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };
double vector[NCOLS] = { 0 };
// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Print vector before Send/Recv
// Create derived contiguous type
MPI_Datatype MY_MPI_ROW;
MPI_Type_contiguous(NCOLS, MPI_DOUBLE, &MY_MPI_ROW);
MPI_Type_commit(&MY_MPI_ROW);
// Send/Recv
tag1 = 1; dest=1; src=0; row=1;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of row %i...\n\n", src, dest, row);
}
if(rank == 0){
    MPI_Send(&matrix[row][0], 1, MY_MPI_ROW,
             dest, tag1, MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&vector[0], NCOLS, MPI_DOUBLE,
             src, tag1, MPI_COMM_WORLD, &Stat);
}
```

Output – 2 Processes

(Rank 0): Matrix:

```
0  1  2  3
4  5  6  7
8  9 10 11
```

(Rank 1): Matrix:

```
0  0  0  0
0  0  0  0
0  0  0  0
```

(Rank 0): Row Vector:

```
0  0  0  0
```

(Rank 1): Row Vector:

```
0  0  0  0
```

Completing Send/Recv from 0 to 1 of
row 1...

(Rank 0): Matrix:

```
0  1  2  3
4  5  6  7
8  9 10 11
```

(Rank 1): Matrix:

```
0  0  0  0
0  0  0  0
0  0  0  0
```

(Rank 0): Row Vector:

```
0  0  0  0
```

(Rank 1): Row Vector:

```
4  5  6  7
```

C – Send One Column from 2D Array (Incorrect)

```
// Incorrect attempt at sending one column of a 2d array
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };

MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
MPI_Comm_rank(MPI_COMM_WORLD,&rank);

// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Send/Recv
tag1 = 1; dest=1; src=0; col=1;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of col %i...\n\n",src,dest,col);
}
if(rank == 0){
    MPI_Send(&matrix[0][col],NROWS,MPI_DOUBLE,
            dest,tag1,MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&matrix[0][col],NROWS,MPI_DOUBLE,
            src,tag1,MPI_COMM_WORLD,&Stat);
}

// Print matrix after Send/Recv
...
```

Output – 2 Processes

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	0	0	0
0	0	0	0
0	0	0	0

Completing Send/Recv from 0 to 1 of
col 1...

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	1	2	3
0	0	0	0
0	0	0	0

Vector Types

- `MPI_Type_vector`: create a type for non-contiguous vectors with constant stride

```
MPI_Type_vector(count,blklen,stride, oldtype,newtype, ierr)
```

ncols

1	6	11	16
2	7	12	17
3	8	13	18
4	9	14	19
5	10	15	20

nrows

count = 4

integer row_type

... cnt blksize stride

```
call MPI_Type_vector(ncols, 1, nrows,  
                    MPI_REAL8, row_type, ierr)
```

```
call MPI_Type_commit(row_type, ierr)
```

Vector Type Example #1

C – Send One Column from 2D Array (MPI_Type_vector)

```
// Sending one column of 2D array via MPI_Type_vector
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };

MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
MPI_Comm_rank(MPI_COMM_WORLD,&rank);

// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Create derived noncontiguous vector type
MPI_Datatype MY_MPI_COLUMN;
MPI_Type_vector(NROWS,1,NCOLS,MPI_DOUBLE,&MY_MPI_COLUMN);
MPI_Type_commit(&MY_MPI_COLUMN);
// Send/Recv
tag1 = 1; dest=1; src=0; col=2;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of col %i...\n\n",src,dest,col);
}
if(rank == 0){
    MPI_Send(&matrix[0][col],1,MY_MPI_COLUMN,
             dest,tag1,MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&matrix[0][col],1,MY_MPI_COLUMN,
             src,tag1,MPI_COMM_WORLD,&Stat);
```

Output – 2 Processes

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	0	0	0
0	0	0	0
0	0	0	0

Completing Send/Recv from 0 to 1 of
col 2...

(Rank 0): Matrix:

0	1	2	3
4	5	6	7
8	9	10	11

(Rank 1): Matrix:

0	0	2	0
0	0	6	0
0	0	10	0

Vector Type Example #2

C – Send One Column from 2D Array to 1D Array

```
// Sending one column of 2D array to 1D array
#define NROWS 3
#define NCOLS 4
double matrix[NROWS][NCOLS] = { 0 };
double vector[NCOLS] = { 0 };
// Initialize matrix on process 0 only
// Print matrix before Send/Recv
// Print vector before Send/Recv
// Send/Recv
tag1 = 1; dest=1; src=0; col=2;
if(rank == 0) {
    printf("\nCompleting Send/Recv from");
    printf(" %i to %i of col %i...\n\n",src,dest,col);
}
MPI_Datatype MY_MPI_COLUMN;
MPI_Type_vector(NROWS,1,NCOLS,MPI_DOUBLE,&MY_MPI_COLUMN);
MPI_Type_commit(&MY_MPI_COLUMN);
if(rank == 0){
    MPI_Send(&matrix[0][col],1,MY_MPI_COLUMN,
            dest,tag1,MPI_COMM_WORLD);
}
if(rank == 1){
    MPI_Recv(&vector[0],NROWS,MPI_DOUBLE,
            src,tag1,MPI_COMM_WORLD,&Stat);
}

// Print matrix after Send/Recv
// Print vector after Send/Recv
```

Output – 2 Processes

(Rank 0): Matrix:

```
0  1  2  3
4  5  6  7
8  9 10 11
```

(Rank 1): Matrix:

```
0  0  0  0
0  0  0  0
0  0  0  0
```

(Rank 0): Col Vector:

0...

(Rank 1): Col Vector:

0...

Completing Send/Recv from 0 to 1 of
col 2...

(Rank 0): Matrix:

```
0  1  2  3
4  5  6  7
8  9 10 11
```

(Rank 1): Matrix:

0...



(Rank 0): Col Vector:

```
0
0
0
```

(Rank 1): Col Vector:

```
2
6
10
```

Derived types (arguments)



Fortran Gotcha!

Fortran – Don't Use Subarray Notation!

```
! Incorrect attempt at sending one row of a 2D array
! Temporary contiguous copy goes out of scope with
    non-blocking calls!
integer, parameter :: NROWS = 3
integer, parameter :: NCOLS = 4
double precision, dimension(NROWS,NCOLS) :: matrix
integer reqs(2)
! Initialize matrix on process 0 only
! Print matrix before Send/Recv
! Send/Recv
tag1 = 1; dest = 1; src = 0; row = 2
if (rank == 0) then ! NO NO NO
    call
        MPI_ISEND(matrix(row,:),NCOLS,MPI_DOUBLE_PRECISION,
        &
        dest,tag1,MPI_COMM_WORLD,reqs(1),ierr)
    call MPI_WAIT(reqs(1),stats(1,1),ierr)
end if
if (rank == 1) then ! NO NO NO
    call
        MPI_Irecv(matrix(row,:),NCOLS,MPI_DOUBLE_PRECISION,
        &
        src,tag1,MPI_COMM_WORLD,reqs(2),ierr)
    call MPI_WAIT(reqs(2),stats(1,2),ierr)
end if
! Print matrix after Send/Recv
...
```

Output – 2 Processes

(Rank 0): Matrix:

```
0.  1.  2.  3.
4.  5.  6.  7.
8.  9. 10. 11.
```

(Rank 1): Matrix:

```
0.  0.  0.  0.
0.  0.  0.  0.
0.  0.  0.  0.
```

Completing Send/Recv from 0 to 1 from
row 2...

(Rank 0): Matrix:

```
0.  1.  2.  3.
4.  5.  6.  7.
8.  9. 10. 11.
```

(Rank 1): Matrix:

```
0.  0.  0.  0.
0.  0.  0.  0.
0.  0.  0.  0.
```

Program received signal SIGSEGV:
Segmentation fault - invalid
memory reference.

Communicators

- A communicator is a “context” for communicating only among a group of tasks.
- `MPI_COMM_WORLD` is the default communicator and consists of all tasks.
- Communication is isolated to context of the group– i.e. no messages from other contexts are “seen”.

Why Communicators?

- **Isolate communication** to a small number of processors
- Useful for creating libraries
- **Collective communication between subgroups** (in lieu of all tasks) can drastically reduce communication costs if only some need to participate
- Useful for communicating with "nearest neighbors"

Communicators and libraries

- Sharing communicator between main and library:
- Library can receive messages from the main program. Oops.

```
main () {  
    if (me==0) MPI_Send( ..to 1.. ,  
                        MPI_COMM_WORLD )  
    library_call()  
    if (me==1) MPI_Recv( .. from 0 .. ,  
                        MPI_COMM_WORLD )  
}  
  
void library_call() {  
    other = me-1;  
    MPI_Recv( .. from other .. ,  
             MPI_COMM_WORLD )  
}
```

Duplicate communicators

- Duplicate communicator with MPI_Comm_dup:
- Same group of processors, but different context: no confusion possible.

```
main () {  
    if (me==0) MPI_Send( ..to 1..,  
                        MPI_COMM_WORLD )  
    library_call(MPI_COMM_WORLD)  
    if (me==1) MPI_Recv( .. from 0 .. ,  
                        MPI_COMM_WORLD )  
}  
  
void library_call(comm) {  
    MPI_Comm my_comm = // copy of comm  
    other = me-1;  
    MPI_Recv( .. from other ..,  
             my_comm )  
}
```

Groups

A new communication group can only be created from a previously defined group. A group must also have a context for communication and, therefore, must have a communicator created for it. The basic steps to form a group are:

- Obtain a complete set of task IDs from a communicator
`MPI_Comm_group`.
- Create a group as a subset of the complete set by
`MPI_Group_excl`, `MPI_Group_incl`, ...
- Create the new communicator for group (subset) using
`MPI_Comm_create`.

Communicators

Routine	Function
<code>MPI_Comm_group</code>	returns group reference of a communicator
<code>MPI_Group_incl</code>	forms new group from inclusion list
<code>MPI_Group_excl</code>	forms new group from exclusion list
<code>MPI_Group_{union, intersection, difference}</code>	Forms new group from union, intersection, or difference of 2 groups.
<code>MPI_Comm_create</code>	creates communicator from a group reference

Communicator Example #1

C Version

```
#define NPROCS 8
int main(int argc, char *argv[]) {
    int rank, new_rank, sendbuf, recvbuf, numprocs;
    int ranks1[4]={0,1,2,3}, ranks2[4]={4,5,6,7};
    MPI_Group orig_group, new_group;
    MPI_Comm new_comm;
    sendbuf = rank;
    // Extract the original group handle
    MPI_Comm_group(MPI_COMM_WORLD, &orig_group);

    // Divide tasks into two distinct groups based upon
    rank
    if (rank < NPROCS/2) {
        MPI_Group_incl(orig_group, NPROCS/2, ranks1,
            &new_group);
    }
    else {
        MPI_Group_incl(orig_group, NPROCS/2, ranks2,
            &new_group);
    }

    // Create new new communicator and then perform
    collective communications
    MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);
    MPI_Allreduce(&sendbuf, &recvbuf, 1, MPI_INT, MPI_SUM,
        new_comm);

    MPI_Group_rank(new_group, &new_rank);
    printf("rank= %d newrank= %d recvbuf=
        %d\n", rank, new_rank, recvbuf);
}
```

Output – 8 Processes

```
rank= 0 newrank= 0 recvbuf= 6
rank= 1 newrank= 1 recvbuf= 6
rank= 2 newrank= 2 recvbuf= 6
rank= 3 newrank= 3 recvbuf= 6
rank= 4 newrank= 0 recvbuf= 22
rank= 5 newrank= 1 recvbuf= 22
rank= 6 newrank= 2 recvbuf= 22
rank= 7 newrank= 3 recvbuf= 22
```


Communicator Example #1

Fortran Version

```
! Divide MPI_COMM_WORLD into two distinct Comms using
! COMM_GROUP, GROUP_INCL, COMM_CREATE
integer, parameter :: NPROCS = 8
integer rank, new_rank, sendbuf, recvbuf, numprocs
integer ranks1(4), ranks2(4), ierr
integer orig_group, new_group, new_comm
data ranks1 /0, 1, 2, 3/, ranks2 /4, 5, 6, 7/
sendbuf = rank
! Extract the original group handle
call MPI_COMM_GROUP(MPI_COMM_WORLD, orig_group, ierr)
! Divide tasks into two distinct groups based upon rank
if (rank .lt. NPROCS/2) then
    call MPI_GROUP_INCL(orig_group, NPROCS/2, &
                        ranks1, new_group, ierr)
else
    call MPI_GROUP_INCL(orig_group, NPROCS/2, &
                        ranks2, new_group, ierr)
endif
call MPI_COMM_CREATE(MPI_COMM_WORLD, new_group, &
                    new_comm, ierr)
call MPI_ALLREDUCE(sendbuf, recvbuf, 1, MPI_INTEGER, &
                  MPI_SUM, new_comm, ierr)
call MPI_GROUP_RANK(new_group, new_rank, ierr)
write(*,*) "(Rank ", rank, "): newrank: ", new_rank, "
           recvbuf: ", recvbuf
```

Output – 8 Processes

```
(Rank 0 ): newrank: 0  recvbuf: 6
(Rank 1 ): newrank: 1  recvbuf: 6
(Rank 2 ): newrank: 2  recvbuf: 6
(Rank 3 ): newrank: 3  recvbuf: 6
(Rank 4 ): newrank: 0  recvbuf: 22
(Rank 5 ): newrank: 1  recvbuf: 22
(Rank 6 ): newrank: 2  recvbuf: 22
(Rank 7 ): newrank: 3  recvbuf: 22
```

MPI_Comm_split

- Provides a short cut method to create a collection of communicators
- All processors with the "same color" will be in the same communicator
- Index controls relative rank in group
- Fortran

```
MPI_Comm_split(OLD_COMM, color, index, NEW_COMM, ierr)
```

- C

```
MPI_Comm_split(OLD_COMM, color, index, &NEW_COMM)
```

Communicator Example #2

C Version

```
#define NPROCS 8
int main(int argc, char *argv[]) {
    int rank, new_rank, sendbuf, recvbuf, numprocs, color;
    MPI_Comm new_comm;

    sendbuf = rank;

    // Divide tasks into two distinct colors based upon
    rank
    if (rank < NPROCS/2) { color = 1; }
    else { color = 2; }

    // Create new new communicator and then perform
    collective communications
    printf("(Rank %i): color: %i\n",rank,color);
    MPI_Comm_split(MPI_COMM_WORLD, color, rank, &new_comm);
    MPI_Allreduce(&sendbuf, &recvbuf, 1, MPI_INT, MPI_SUM,
        new_comm);

    MPI_Comm_rank(new_comm, &new_rank);
    printf("(Rank %i): newrank: %i recvbuf:
        %i\n",rank,new_rank,recvbuf);
```

Output – 8 Processes

```
(Rank 1): color: 1
(Rank 2): color: 1
(Rank 3): color: 1
(Rank 4): color: 2
(Rank 5): color: 2
(Rank 6): color: 2
(Rank 7): color: 2
(Rank 0): color: 1
(Rank 0): newrank: 0 recvbuf: 6
(Rank 1): newrank: 1 recvbuf: 6
(Rank 2): newrank: 2 recvbuf: 6
(Rank 3): newrank: 3 recvbuf: 6
(Rank 4): newrank: 0 recvbuf: 22
(Rank 5): newrank: 1 recvbuf: 22
(Rank 6): newrank: 2 recvbuf: 22
(Rank 7): newrank: 3 recvbuf: 22
```

Communicator Example #2

Fortran Version

```
! Divide MPI_COMM_WORLD into two distinct Comms using
! COMM_GROUP, GROUP_INCL, COMM_CREATE
integer, parameter :: NPROCS = 8
integer rank, new_rank, sendbuf, recvbuf, numprocs
integer color, new_comm, ierr

sendbuf = rank

! Divide tasks into two distinct groups based upon rank
if (rank .lt. NPROCS/2) then
    color = 1
else
    color = 2
endif

write(*,*)"Rank ",rank,": color: ",color
call MPI_COMM_SPLIT(MPI_COMM_WORLD, color, rank,
    new_comm, ierr)
call MPI_ALLREDUCE(sendbuf,recvbuf,1,MPI_INTEGER, &
    MPI_SUM, new_comm, ierr)

call MPI_COMM_RANK(new_comm, new_rank, ierr)
write(*,*)"Rank ",rank,": newrank: ",new_rank,"
    recvbuf: ",recvbuf
```

Output – 8 Processes

```
(Rank 1): color: 1
(Rank 2): color: 1
(Rank 3): color: 1
(Rank 5): color: 2
(Rank 6): color: 2
(Rank 7): color: 2
(Rank 0): color: 1
(Rank 4): color: 2
(Rank 0): newrank: 0  recvbuf: 6
(Rank 1): newrank: 1  recvbuf: 6
(Rank 2): newrank: 2  recvbuf: 6
(Rank 3): newrank: 3  recvbuf: 6
(Rank 4): newrank: 0  recvbuf: 22
(Rank 5): newrank: 1  recvbuf: 22
(Rank 7): newrank: 3  recvbuf: 22
(Rank 6): newrank: 2  recvbuf: 22
```

Topologies

- Use the MPI library for common grid topologies (**local functions**)
- A *topology* maps process-ranks onto a set of N-tuples.
- E.g. {0, 1, 2, 3} → {(0,0), (0,1), (1,0), (1,1)} (**row-major in ranks**)
- Cartesian Maps (**arbitrary number of dimensions**):
 - MPI_Cart_create Creates map (ranks → coordinates).
 - MPI_Cart_get Returns info created in MPI_Cart_create.
 - MPI_Cart_coords Returns coordinates from rank.
 - MPI_Cart_rank Returns rank from coordinates.
 - MPI_Cart_shift Returns Nth neighbor's coords.
- **graph** constructors go beyond the *N*-dimensional rectilinear mapping of the Cartesian topology (**MPI_Graph_create**)

Note: the virtual topology does not necessarily map the hardware processor grid to the process grid in the most efficient manner.

:

(Virtual) Topologies

- In terms of MPI, a virtual topology **describes a mapping and ordering of MPI processes into a geometric shape.**
- The two main types of topology supported by MPI are **Cartesian(grid) and Graph.**
- MPI topologies are **virtual – there may be no relation between the physical structure** of parallel machine and the process topology.
- Virtual topologies are **built upon MPI communicator and groups.**
- Must be ***programmed* by the application developer.**
- Useful for applications with **specific communication pattern.**
- A particular **implementation may optimize process mapping** based on the physical characteristics of a given parallel machine.
- Can be used within an intra-communicator; cannot be added to inter-communicators.

References

- Victor Eijkhout, “Introduction to High Performance Scientific Computing”
- Victor Eijkhout, “Parallel Computing for Science and Engineering”
- Mark Lubin, “Introduction into new features of MPI-3.0 Standard”
- “MPI: A Message-Passing Interface Standard Version-3.0”