

PCSE Lecture 13

HPC Linear Algebra

Cyrus Proctor
Victor Eijkhout

April 23, 2015



Review

- Start out with a partial differential equation (BVP, IBVP, etc.)
- Discretize in space and time
 - Explicit
 - Implicit
- Von Neumann stability analysis
- Explicit: compute matrix vector products (MVPs)
- Implicit: compute linear system solution
- Linear system solutions
 - Direct methods
 - Iterative methods
- Iterative methods
 - Use preconditioner K to converge faster
 - MVPs, vector addition, vector inner products common
- Sparse matrix storage


Sparse Matrix Vector Products

Recall the two storage types mentioned so far:

- Diagonal storage
- Compressed row storage



Sparse Matrix Vector Products



Data reuse in the SMVP:

- Similarities between dense MVP by row and sparse MVP:
 - All matrix elements are used sequentially
 - Any cache line loaded is utilized fully
- Difference for CRS MVP:
 - Indirect addressing requires loading the elements of an integer vector
 - More memory traffic for the same number of operations
 - Elements of the source vector are **not** loaded sequentially
 - Cacheline with source element likely to be not fully utilized

Codes dominated by sparse MVPs may only run at $\sim 5\%$ peak


Vectorization of the Sparse MVP

In some circumstances, bandwidth and reuse are not the dominant concerns:

- Because of all the zero elements, CRS vector lengths are typically low
- GPUs/MICs expect identical operations for best efficiency
- Sparse rows in CRS are likely unequal length
- Efforts to use variations of diagonal storage have seen a revival lately

How to Improve Performance?

- If there is any regular pattern to non-zeros
- Small dense blocks
- Reduces indexing information
 - Just using 2x2 blocks gives 4X reduction in index data needed



Ordering Strategies and Parallelism

For the following methods, keep in mind that we are:

- Renumbering grid points
- Redistributing which process(es) have access to data

These strategies work well if:

- Recursion length is “large enough”
 - Dense systems usually okay
 - More work for sparse systems

All strategies are variants of Gaussian elimination.

Variable Reordering

Consider our 1D BVP tridiagonal matrix with **Lexicographically** ordered points:

$$\begin{pmatrix} a_{11} & a_{12} & & & \emptyset \\ a_{21} & a_{22} & a_{23} & & \\ & a_{32} & a_{33} & a_{34} & \\ \emptyset & & \ddots & \ddots & \ddots \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \vdots \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ \vdots \end{pmatrix}$$

- Matrix A is known
- Vector b is known
- Vector x is **unknown**

How to go about maximizing available parallelism?

Red-Black Ordering

- x_j **only** depends on x_{j-1} and x_{j+1}

$$\begin{pmatrix} a_{11} & & & & a_{12} & & & \\ & a_{33} & & & a_{32} & a_{34} & & \\ & & a_{55} & & & \ddots & \ddots & \\ & & & \ddots & & & & \\ a_{21} & a_{23} & & & a_{22} & & & \\ & a_{43} & a_{45} & & & a_{44} & & \\ & & \ddots & \ddots & & & \ddots & \end{pmatrix} \begin{pmatrix} x_1 \\ x_3 \\ x_5 \\ \vdots \\ x_2 \\ x_4 \\ \vdots \end{pmatrix} = \begin{pmatrix} b_1 \\ b_3 \\ b_5 \\ \vdots \\ b_2 \\ b_4 \\ \vdots \end{pmatrix}$$



Red-Black Ordering 2

Using this grouping and denoting with subscripts (b) and (r) the black and red nodes, respectively, we can write for the exact system:

$$\begin{pmatrix} D_r & U \\ L & D_b \end{pmatrix} \begin{pmatrix} x_r \\ x_b \end{pmatrix} = \begin{pmatrix} b_r \\ b_b \end{pmatrix}$$

We can turn this into an iterative update:

$$\begin{pmatrix} D_r & U \\ L & D_b \end{pmatrix} \begin{pmatrix} x_r^{k+1} \\ x_b^{k+1} \end{pmatrix} = \begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix} \begin{pmatrix} x_r^k \\ x_b^k \end{pmatrix} + \begin{pmatrix} b_r \\ b_b \end{pmatrix}$$

Choice of $K = A$, solution in “1 iteration”, i.e. direct solution

Red-Black Ordering 3

Instead, move L and U to right-hand side (old-iterate) for $K = D_A$ Jacobi preconditioner:

$$\begin{pmatrix} D_r & 0 \\ 0 & D_b \end{pmatrix} \begin{pmatrix} x_r^{k+1} \\ x_b^{k+1} \end{pmatrix} = \begin{pmatrix} 0 & -U \\ -L & 0 \end{pmatrix} \begin{pmatrix} x_r^k \\ x_b^k \end{pmatrix} + \begin{pmatrix} b_r \\ b_b \end{pmatrix}$$

Use (sparse) MVPs to solve:

$$D_r x_r^{k+1} = -U x_b^k + b_r$$

$$D_b x_b^{k+1} = -L x_r^k + b_b$$

Which becomes:

$$x_r^{k+1} = D_r^{-1}[-U x_b^k + b_r]$$

$$x_b^{k+1} = D_b^{-1}[-L x_r^k + b_b]$$

So, we've only **changed the order** that we solve the x 's

Red-Black Ordering 4

Now, only move U to the right-hand side for

$K = D_A + L_A$ Gauss-Seidel preconditioner:

$$\begin{pmatrix} D_r & 0 \\ L & D_b \end{pmatrix} \begin{pmatrix} x_r^{k+1} \\ x_b^{k+1} \end{pmatrix} = \begin{pmatrix} 0 & -U \\ 0 & 0 \end{pmatrix} \begin{pmatrix} x_r^k \\ x_b^k \end{pmatrix} + \begin{pmatrix} b_r \\ b_b \end{pmatrix}$$

Use (sparse) MVPs to solve:

$$D_r x_r^{k+1} = -U x_b^k + b_r$$

$$L x_r^{k+1} + D_b x_b^{k+1} = b_b$$

Which becomes (in parallel):


Step 1: $x_r^{k+1} = D_r^{-1}[-U x_b^k + b_r]$

Step 2: Distribute relevant parts of x_r^{k+1}

Step 3: $x_b^{k+1} = D_b^{-1}[-L x_r^{k+1} + b_b]$

Red-Black Ordering 5


In serial:



“Distribution” step is unnecessary. All points of x_r^{k+1} are available locally.

Red-Black Ordering 6


In parallel:



Distribute points of x_r^{k+1} that are needed by neighbor processes.

Red-Black Ordering 7

- Red-black ordering can be applied in 2 dimensions just as well
- Apply to points (i, j) where $1 \leq i, j \leq n$
- Successive odd points $(1, 1), (1, 3), (1, 5) \dots$ on the first line
- Successive even points $(2, 2), (2, 4), (2, 6) \dots$ on the second line
- Successive odd points ... on the third line




Setup is otherwise the same!

Nested Dissection Ordering A.K.A Domain Decomposition

- Initially designed as a way to reduce fill-in
- Advantageous in a parallel computing context
- Recursive process

First, split computational domain into two parts (1, 2) with a
dividing strip (3):



Nested Dissection 2

- The divider must be wide enough such that the other two subdomains are not connected, i.e. decoupled.

Our resulting matrix, A^{DD} has the following structure:

$$\left(\begin{array}{ccccc|ccccc|c}
 \star & \star & & & & & & & & & 0 \\
 \star & \star & \star & & & & & & & & \vdots \\
 & & \ddots & \ddots & \ddots & & & & & & \vdots \\
 & & & \star & \star & \star & & & & & 0 \\
 & & & & \star & \star & & & & & \star \\
 \hline
 & & & & & \star & \star & & & & 0 \\
 & & & & & \star & \star & \star & & & \vdots \\
 & & & & & & \ddots & \ddots & \ddots & & \vdots \\
 & & & & & & & \star & \star & \star & 0 \\
 & & & & & & & & \star & \star & \star \\
 \hline
 0 & \dots & \dots & 0 & \star & 0 & \dots & \dots & 0 & \star & \star
 \end{array} \right) \left. \begin{array}{l} \\ \\ \\ \\ \\ \\ \\ \\ \\ \\ \end{array} \right\} \begin{array}{l} \\ \\ \\ (n^2 - n)/2 \\ \star \\ \\ \\ \\ (n^2 - n)/2 \\ n \end{array}$$

Nested Dissection 3

In block matrix form:

$$A^{DD} = \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ \emptyset & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{pmatrix}$$

- This gives a 3×3 block matrix structure
- Corresponds to the 3 divisions of the computational domain
- A_{12} and A_{21} are zero b/c subdomains 1 & 2 are decoupled
- Subdomain 3 contains coupling terms
- Subdomain 3 is the **parent** of subdomains 1 & 2
- Subdomains 1 & 2 are **siblings**

Nested Dissection 4

Factorize the 3×3 block matrix using an LU decomposition:


$$A^{DD} = LU = \begin{pmatrix} I & & \\ \emptyset & I & \\ A_{31}A_{11}^{-1} & A_{32}A_{22}^{-1} & I \end{pmatrix} \begin{pmatrix} A_{11} & \emptyset & A_{13} \\ & A_{22} & A_{23} \\ & & S_{33} \end{pmatrix}$$

- Where $S_{33} = A_{33} - A_{31}A_{11}^{-1}A_{13} - A_{32}A_{22}^{-1}A_{23}$
- Can compute $A_{31}A_{11}^{-1}A_{13}$ and $A_{32}A_{22}^{-1}A_{23}$ **simultaneously!**
- The third subdomain remains **sequential**

Nested Dissection 5

Let's start over:

- With one division, we got a factorization that was two-way parallel
- Instead of one vertical division, try one vertical division and one horizontal division



Nested Dissection 6

This gives a block matrix form:

$$A^{DD} = \left(\begin{array}{cccc|cc|c} A_{11} & & & & A_{15} & & A_{17} \\ & A_{22} & & & A_{25} & & A_{27} \\ & & A_{33} & & & A_{36} & A_{37} \\ & & & A_{44} & & A_{46} & A_{47} \\ \hline A_{51} & A_{52} & & & A_{55} & & A_{57} \\ & & A_{63} & A_{64} & & A_{66} & A_{67} \\ \hline A_{71} & A_{72} & A_{73} & A_{74} & A_{75} & A_{76} & A_{77} \end{array} \right)$$

- Subdomains 1, 2, 3, and 4 are siblings
- Subdomains 5 and 6 are parents
- Subdomain 7 is a grandparent

Nested Dissection 7

The U from the LU factorization of A^{DD} is:

$$U = \left(\begin{array}{ccc|cc|c} A_{11} & & & A_{15} & & A_{17} \\ & A_{22} & & A_{25} & & A_{27} \\ & & A_{33} & & A_{36} & A_{37} \\ & & & A_{44} & A_{46} & A_{47} \\ \hline & & & & S_5 & A_{57} \\ & & & & & S_6 \\ \hline & & & & & S_7 \end{array} \right)$$

- Siblings are computed first, then parents, then grandparents
 - ① Compute simultaneously sibling factorizations (more DD's?!)
 - ② When siblings are finished, parents start (still parallel components)
 - ③ When parents finish, grandparent starts (sequential)

Nested Dissection 8

- Above process may complete until subdomains are very small
- Theoretically can go to 1×1
- Usually go to $\sim 32 \times 32$ blocks and use efficient dense solver
- Tend to see nice space savings compared to regular factorization with general 2D case
- 3D tends to have higher complexity

Nested Dissection 9

- Mapping these tasks to processors is not trivial; use task queue
- Will end up with more processors than tasks towards the end
- Last tasks are most substantial! Get free processors to help

Task Queue

For(all bottom level subdomains d)

 add d to the Queue

While(Queue is not empty)

If(a processor is idle)(assign a queued task to it)

If(a task is finished AND its sibling is finished)
 add its parent to the queue

- For distributed memory, communication is done intelligently
- Still relatively cheap compared to factorization

Grid Updates


- Sparse MVPs are largely **bandwidth-bound**
- Little data-reuse; indexing make locality worse
- Try to rearrange operations to lessen bandwidth demands

Consider our 1D IBVP Explicit scheme:


$$\forall j: x_j^{k+1} = f(x_j^k, x_{j-1}^k, x_{j+1}^k)$$

- Assume the set $\{x_j^k\}_j$ is too large to fit in cache

Schematically:




Grid Updates 2




- Typically compute all x_j^{k+1} values
- Then, compute all x_j^{k+2} values using x_j^{k+1} 's
- x_j^{k+1} values are flushed from cache as they are generated
- Why not **compute**, not one, but **two k iterations** for a given j ?
- x_0^{k+2} requires x_0^{k+1} and x_1^{k+1}
 - x_0^{k+1} requires x_0^k and x_1^k
 - x_1^{k+1} requires x_0^k , x_1^k , and x_2^k

Grid Updates 3



- Suppose we only care about the final result
- Processor 0 (red) computes 4 “ j ” points on “ k ” level ($k + 2$)
- For this, it needs 5 j points from level ($k + 1$)
- These points need to be computed from 6 j points from level (k)
- Processor 0 needs a “ghost region” (or halo) of j size 2 instead of 1 for a single k update

Grid Updates 4



- One of the points computed by processor 1 is x_4^{k+2}
- x_4^{k+2} needs x_4^{k+1}
- x_3^{k+2} also needs x_4^{k+1}

Which processor should compute x_4^{k+1} ?

- Why not both processor 0 and processor 1?

Grid Updates 5


But you're doing more work?!

- Do a colored block on a single processor
- If all points in a colored block fit in cache, all $(k + 1)$ points are reused before flushing

For distributed memory computation:

- Reduces message traffic
- Neighbor processors exchange every $k = 2$ steps instead of every iteration step
- Redundant calculations in this arrangement are typically faster than an extra communication

Grid Updates 6



Communication and work minimizing strategy:

- Make algorithm more efficient
- Processor i communicates its halo points at step (k) to neighbor processors
- Neighbor processors simultaneously do the same
- Processor i , **does not wait** to start $(k + 1)$ and then $(k + 2)$ updates that already can be done locally
- Once all updates that can be done locally have finished, block for halo communication
- Compute missing values in halo

References

- Victor Eijkhout, “Introduction to High Performance Scientific Computing”
- Victor Eijkhout, “Parallel Computing for Science and Engineering”
- George Karniadakis “Parallel Scientific Computing in C++ and MPI”