

# Introduction to PETSc

Victor Eijkhout

# Outline

- 1 Application context
- 2 Library introduction
- 3 Vector and matrix objects
- 4 Iterative methods
- 5 Tools and such

# Table of Contents

- 1 Application context
- 2 Library introduction
- 3 Vector and matrix objects
- 4 Iterative methods
- 5 Tools and such

# Sparse matrix from 2D PDE

Two-dimensional:  $-u_{xx} - u_{yy} = f$  on unit square  $[0, 1]^2$

Difference equation:


$$4u(x, y) - u(x + h, y) - u(x - h, y) - u(x, y + h) - u(x, y - h) = h^2 f(x, y)$$

$$4u_k - u_{k-1} - u_{k+1} - u_{k-n} - u_{k+n} = f_k$$

Consider a graph where  $\{u_k\}_k$  are the edges  
and  $(u_i, u_j)$  is an edge iff  $a_{ij} \neq 0$ .

# The graph view of things

Poisson eq:



This is a graph!

This is the (adjacency) graph of a sparse matrix.

# The matrix view of things

$$\left( \begin{array}{cccc|cccc|c}
 4 & -1 & & & 0 & -1 & & & 0 \\
 -1 & 4 & 1 & & & & -1 & & \\
 & \ddots & \ddots & \ddots & & & & \ddots & \\
 & & \ddots & \ddots & -1 & & & & \\
 0 & & & -1 & 4 & 0 & & & -1 \\
 \hline
 -1 & & & & 0 & 4 & -1 & & -1 \\
 & -1 & & & & -1 & 4 & -1 & \\
 & \uparrow & \ddots & & & \uparrow & \uparrow & \uparrow & \\
 & k-n & & & & k-1 & k & k+1 & -1 \\
 & & & & -1 & & & -1 & 4 \\
 \hline
 & & & & & \ddots & & & \ddots
 \end{array} \right)$$

# LU of a sparse matrix

$$\Rightarrow \left( \begin{array}{ccccc|ccc} 4 & -1 & 0 & \dots & & -1 & & \\ -1 & 4 & -1 & 0 & \dots & 0 & -1 & \\ & \ddots & \ddots & \ddots & & & \ddots & \\ - & - & - & - & - & - & - & - \\ -1 & 0 & \dots & & & 4 & -1 & \\ 0 & -1 & 0 & \dots & & -1 & 4 & -1 \end{array} \right)$$

$$\Rightarrow \left( \begin{array}{c|cccc|ccc} 4 & -1 & 0 & \dots & & -1 & & \\ \hline & 4 - \frac{1}{4} & -1 & 0 & \dots & -1/4 & -1 & \\ & \ddots & \ddots & \ddots & & & \ddots & \\ - & - & - & - & - & - & - & - \\ & -1/4 & \dots & & & 4 - \frac{1}{4} & -1 & \\ & -1 & 0 & \dots & & -1 & 4 & -1 \end{array} \right)$$

## Fill-in during LU

2D BVP:  $\Omega$  is  $n \times n$ , gives matrix of size  $N = n^2$ , with bandwidth  $n$ .

Matrix storage  $O(N)$

LU storage  $O(N^{3/2})$  (limited to band)

LU factorization work  $O(N^2)$

Complicated recurrences hard to parallelize



# Table of Contents

- 1 Application context
- 2 Library introduction**
- 3 Vector and matrix objects
- 4 Iterative methods
- 5 Tools and such

# What is PETSc? Why should you use PETSc?

Portable Extendable Toolkit for Scientific Computations


- Scientific Computations: parallel linear algebra, in particular linear and nonlinear solvers
- Toolkit: Contains high level solvers, but also the low level tools to roll your own.
- Portable: Available on many platforms, basically anything that has MPI

Why use it? It's big, powerful, well supported.

# What is in PETSc?

- Linear system solvers (sparse/dense, iterative/direct)
- Nonlinear system solvers
- Tools for distributed matrices
- Support for profiling, debugging, graphical output

Level of  
Abstraction



Application Codes



SNES  
(Nonlinear Equations Solvers)

TS  
(Time Stepping)

PC  
(Preconditioners)

KSP  
(Krylov Subspace Methods)

Matrices

Vectors

Index Sets

BLAS

MPI

# Parallel Numerical Components of PETSc

Nonlinear Solvers		
Newton-based Methods		Other
Line Search	Trust Region	

Time Steppers			
Euler	Backward Euler	Pseudo-Time Stepping	Other

Krylov Subspace Methods							
GMRES	CG	CGS	Bi-CG-Stab	TFQMR	Richardson	Chebyshev	Other

Preconditioners						
Additive Schwarz	Block Jacobi	Jacobi	ILU	ICC	LU (sequential only)	Other

Matrices				
Compressed Sparse Row (AIJ)	Block Compressed Sparse Row (BAIJ)	Block Diagonal (BDiag)	Dense	Other

<b>Vectors</b>
----------------

Index Sets			
Indices	Block Indices	Stride	Other

# PETSc and parallelism

PETSc is layered on top of MPI

MPI has basic tools: send elementary datatypes between processors

PETSc has intermediate tools:

insert matrix element in arbitrary location,  
do parallel matrix-vector product

⇒ you do not need to know much MPI when you use PETSc

# Table of Contents


- 1 Application context
- 2 Library introduction
- 3 Vector and matrix objects**
- 4 Iterative methods
- 5 Tools and such

# Parallel layout

Local or global size in


```
VecSetSizes(Vec v, int m, int M);
```

Global size can be specified as `PETSC_DECIDE`.



`VecSetSizes(V,2,5)`

`VecSetSizes(V,3,5)`



`VecSetSizes(V,2,PETSC_DECIDE)`


`VecSetSizes(V,3,PETSC_DECIDE)`



# Parallel layout up to PETSc

```
VecSetSizes(Vec v, int m, int M);
```

Local size can be specified as `PETSC_DECIDE`.



```
VecSetSizes(V,PETSC_DECIDE,8)
```

```
VecSetSizes(V,PETSC_DECIDE,8)
```

```
VecSetSizes(V,PETSC_DECIDE,8)
```

# Setting values

Set vector to constant value:

```
VecSet(Vec x,PetscScalar value);
```

Set individual elements (global indexing!):

```
VecSetValues(Vec x,int n,int *indices,PetscScalar *values,  
             INSERT_VALUES); /* or ADD_VALUES */
```

```
i = 1; v = 3.14;  
VecSetValues(x,1,&i,&v,INSERT_VALUES);  
ii[0] = 1; ii[1] = 2; vv[0] = 2.7; vv[1] = 3.1;  
VecSetValues(x,2,ii,vv,INSERT_VALUES);
```

```
call VecSetValues(x,1,i,v,INSERT_VALUES,ierr,e)  
call VecSetValues(x,2,ii,vv,INSERT_VALUES,ierr,e)
```

# Setting values

No restrictions on parallelism;  
after setting, move values to appropriate processor:

```
VecAssemblyBegin(Vec x);  
VecAssemblyEnd(Vec x);
```

# Basic operations

```
VecAXPY(Vec y,PetscScalar a,Vec x);    /*  $y \leftarrow y + a x$  */
VecAYPX(Vec y,PetscScalar a,Vec x);    /*  $y \leftarrow a y + x$  */
VecScale(Vec x, PetscScalar a);
VecDot(Vec x, Vec y, PetscScalar *r); /* several variants */
VecMDot(Vec x,int n,Vec y[],PetscScalar *r);
VecNorm(Vec x, NormType type, double *r);
VecSum(Vec x, PetscScalar *r);
VecCopy(Vec x, Vec y);
VecSwap(Vec x, Vec y);
VecPointwiseMult(Vec w,Vec x,Vec y);
VecPointwiseDivide(Vec w,Vec x,Vec y);
VecMAXPY(Vec y,int n, PetscScalar *a, Vec x[]);
VecMax(Vec x, int *idx, double *r);
VecMin(Vec x, int *idx, double *r);
VecAbs(Vec x);
VecReciprocal(Vec x);
VecShift(Vec x,PetscScalar s);
```

**Mat Datatype: matrix**

# Matrix creation

The usual create/destroy calls:

```
MatCreate(MPI_Comm comm, Mat *A)
MatDestroy(Mat A)
```

Several more aspects to creation:

```
MatSetType(A, MATSEQAIJ) /* or MATMPIAIJ or MATAIJ */
MatSetSizes(Mat A, int m, int n, int M, int N)
MatSeqAIJSetPreallocation /* more about this later*/
(Mat B, PetscInt nz, const PetscInt nnz[])
```

Local or global size can be PETSC\_DECIDE (as in the vector case)

# Matrix Preallocation

- PETSc matrix creation is very flexible:
- No preset sparsity pattern
- any processor can set any element  
⇒ potential for lots of malloc calls
- (run your code with `-memory_info`, `-malloc_log`)

malloc is very expensive:

tell PETSc the matrix' sparsity structure

(do construction loop twice: once counting, once making)

# Sequential matrix structure


```
MatSeqAIJSetPreallocation
```

```
(Mat B, PetscInt nz, const PetscInt nnz[])
```

- `nz` number of nonzeros per row  
(or slight overestimate)
- `nnz` array of row lengths (or overestimate)
- considerable savings over dynamic allocation!



# Parallel matrix structure



# Parallel matrix structure description

```
MatMPIAIJSetPreallocation(Mat B,  
    PetscInt d_nz,const PetscInt d_nnz[],  
    PetscInt o_nz,const PetscInt o_nnz[])
```

- `d_nz`: number of nonzeros per row in diagonal part
- `o_nz`: number of nonzeros per row in off-diagonal part
- `d_nnz`: array of numbers of nonzeros per row in diagonal part
- `o_nnz`: array of numbers of nonzeros per row in off-diagonal part



In Fortran use `PETSC_NULL_INTEGER` if not specifying arrays

# Setting matrix elements

► Recall laplace graph

```
MatGetOwnershipRange(A, &low, &high);
for ( i=0; i<m; i++ ) {
    for ( j=0; j<n; j++ ) {
        I = j + n*i;
        if (I>=low && I<high) { // my row!
            J = I-1; v = -1.0;
            if (i>0) MatSetValues
                (A, 1, &I, 1, &J, &v, INSERT_VALUES);
            J = I+1 // et cetera
        }
    }
}
MatAssemblyBegin(A, MAT_FINAL_ASSEMBLY);
MatAssemblyEnd(A, MAT_FINAL_ASSEMBLY);
```

# Finite Element Matrix assembly



# Finite Element Matrix assembly

```
for (e=myfirstelement; e<mylastelement; e++) {  
    for (i=0; i<nlocalnodes; i++) {  
        I = localtoglobal(e,i);  
        for (j=0; j<nlocalnodes; j++) {  
            J = localtoglobal(e,j);  
            v = integration(e,i,j);  
            MatSetValues  
                (mat,1,&I,1,&J,&v,ADD_VALUES);  
            ....  
        }  
    }  
}  
  
MatAssemblyBegin(mat,MAT_FINAL_ASSEMBLY);  
MatAssemblyEnd(mat,MAT_FINAL_ASSEMBLY);
```

# Matrix usage


Very important:

```
MatMult(Mat mat,Vec x,Vec y)
```

```
MatMultTranspose(Mat mat,Vec x,Vec y)
```

```
MatMultAdd(Mat mat,Vec v1,Vec v2,Vec v3)
```

Sparse matrix vector product induces halo region;



Product  $y \leftarrow Ax$  becomes

$$y \leftarrow A_{\text{local}}x_{\text{local}} + A_{\text{remote}}x_{\text{remote}}$$

► Recall matrix structure

- Start communication for halo
- Do local computation
- Wait for completion of halo transfer
- Do remote computation

# Querying parallel structure

Matrix partitioned by block rows:

```
MatGetSize(Mat mat,PetscInt *m,PetscInt* n);  
MatGetLocalSize(Mat mat,PetscInt *m,PetscInt* n);  
MatGetOwnershipRange(Mat A,int *first row,int *last row);
```



# Table of Contents

- 1 Application context
- 2 Library introduction
- 3 Vector and matrix objects
- 4 Iterative methods**
- 5 Tools and such

# What are iterative solvers?

Solving a linear system  $Ax = b$  with Gaussian elimination can take lots of time/memory.

Alternative: iterative solvers use successive approximations of the solution:

- Convergence not always guaranteed
- Possibly much faster / less memory
- Basic operation:  $y \leftarrow Ax$  executed once per iteration
- Also needed: preconditioner  $B \approx A^{-1}$

# General form of iterative methods

$$\begin{cases} r_i = Ax_i - b \\ x_{i+1}\gamma_{i+1,i} = K^{-1}r_i + \sum_{j \leq i} x_j \gamma_{ji} \\ r_{i+1}\gamma_{i+1,i} = AK^{-1}r_i + \sum_{j \leq i} r_j \gamma_{ji} \end{cases} \quad \text{where } \gamma_{i+1,i} = \sum_{j \leq i} \gamma_{ji}.$$

# Conjugate Gradients

Basic idea:

$$r_i^t K^{-1} r_j = 0 \quad \text{if } i \neq j.$$

Split recurrences:

$$\begin{cases} x_{i+1} = x_i - \delta_i p_i \\ r_{i+1} = r_i - \delta_i A p_i \\ p_i = K^{-1} r_i + \gamma_i p_{i-1}, \end{cases}$$

where

$$\delta_i = \frac{r_i^t K^{-1} r_i}{p_i^t A p_i}, \quad \gamma_i = \frac{r_i^t K^{-1} r_i}{r_{i-1}^t K^{-1} r_{i-1}}$$

# Components of every iterative method

- Matrix-vector product
- Preconditioner (construction and application)
- Inner products
- Other vector operations.

# Basic concepts

- All linear solvers in PETSc are iterative (see below)
- Object oriented: solvers only need matrix action, so can handle shell matrices
- Preconditioners
- Fargoin control through commandline options
- Tolerances, convergence and divergence reason
- Custom monitors and convergence tests

## KSP: Krylov Space objects

# Iterative solver basics

```
KSPCreate(comm,&solver);  
KSPSetOperators(solver,A,B);  
/* optional */ KSPSetup(solver);  
KSPSolve(solver,rhs,sol);  
KSPDestroy(solver);
```



# Settings in general

- Other settings, both as command and runtime option
- option `-ksp_view`
- (preconditioners discussed below)

# Solver type and tolerances

```
KSPSetType(solver, KSPGMRES);  
KSPSetTolerances(solver, rtol, atol, dtol, maxit);
```

KSP can be controlled from the commandline:

```
KSPSetFromOptions(solver);  
/* right before KSPSolve or KSPSetUp */
```

then options `-ksp....` are parsed.

- **type:** `-ksp_type gmres -ksp_gmres_restart 20`
- **tolerances:** `-ksp_max_it 50`

# Convergence

Iterative solvers can fail

- Solve call itself gives no feedback: solution may be completely wrong
- `KSPGetConvergedReason(solver, &reason)` :  
positive is convergence, negative divergence  
(`${PETSC_DIR}/include/petscksp.h` for list)
- `KSPGetIterationNumber(solver, &nits)` : after how many iterations did the method stop?

## **PC: Preconditioner objects**

# PC basics

- PC usually created as part of KSP: separate create and destroy calls exist, but are (almost) never needed

```
KSP solver; PC precon;  
KSPCreate(comm, &solver);  
KSPGetPC(solver, &precon);  
PCSetType(precon, PCJACOBI);
```

- PCJACOBI, PCILU (only sequential), PCASM, PCBJACOBI, PCMG, et cetera
- Controllable through commandline options:  
`-pc_type ilu -pc_factor_levels 3`

# Direct methods

- Iterative method with direct solver as preconditioner would converge in one step
- Direct methods in PETSc implemented as special iterative method:  
KSPPREONLY only apply preconditioner
- All direct methods are preconditioner type PCLU:

```
myprog -pc_type lu -ksp_type preonly \  
-pc_factor_mat_solve_package mumps
```

# Table of Contents

- 1 Application context
- 2 Library introduction
- 3 Vector and matrix objects
- 4 Iterative methods
- 5 Tools and such**

# Running

Parallel invocation. On your own machine:

```
mpirun -np 3 petscprog <bunch of runtime options>
```

On stampede and such: using `ibrun`

Petsc has lots of runtime options.

- `-log_summary` : give runtime statistics
- `-malloc_dump -memory_info` : memory statistics
- `-start_in_debugger` : parallel debugging (not on our clusters, but very convenient on your laptop)
- `-options_left` : check for mistyped options
- `-ksp_type gmres (et cetera)` : program control



# Library setup, C

```
ierr = PetscInitialize(&argc,&argv,0,0); CHKERRQ(ierr);  
// all the petsc work  
ierr = PetscFinalize();CHKERRQ(ierr);
```

Can replace `MPI_Init`

General: Every routine has an error return. Catch that value!

# Library setup, F

```
call PetscInitialize(PETSC_NULL_CHARACTER,ierr)
CHKERRQ(ierr)
// all the petsc work
call PetscFinalize(ierr)
CHKERRQ(ierr)
```

Error code is now final parameter. This holds for every PETSc routine

## Note to self

```
PetscInitialize  
  (&argc,&args,0,"Usage: prog -o1 v1 -o2 v2\n");
```

run as

```
./program -help
```

This displays the usage note, plus all available petsc options.

Not available in Fortran

# Commandline options, C

```
ierr = PetscOptionsGetInt  
      (PETSC_NULL, "-n", &n, PETSC_NULL); CHKERRQ(ierr);  
ierr = PetscPrintf  
      (comm, "Input parameter: %d\n", n); CHKERRQ(ierr);
```

Read commandline argument, print out from processor zero

## Commandline options, F

```
character*80      msg
call PetscOptionsGetInt(PETSC_NULL_CHARACTER,
>      "-n",n,PETSC_NULL_CHARACTER,ierr)
CHKERRQ(ierr)
write(msg,10) n
10  format("Input parameter:",i5)
call PetscPrintf(PETSC_COMM_WORLD,msg,ierr)
CHKERRQ(ierr)
```

**Note the PETSC\_NULL\_CHARACTER, note that PetscPrintf has only one string argument**

## Profiling, debugging

# Basic profiling

- `-log_summary` flop counts and timings of all PETSc events
- `-info` all sorts of information, in particular

```
%% mpiexec yourprogram -info | grep malloc
[0] MatAssemblyEnd_SeqAIJ():
    Number of mallocs during MatSetValues() is 0
```
- `-log_trace` start and end of all events: good for hanging code

## Log summary: overall

	Max	Max/Min	Avg	Total
Time (sec):	5.493e-01	1.00006	5.493e-01	
Objects:	2.900e+01	1.00000	2.900e+01	
Flops:	1.373e+07	1.00000	1.373e+07	2.746e+07
Flops/sec:	2.499e+07	1.00006	2.499e+07	4.998e+07
Memory:	1.936e+06	1.00000		3.871e+06
MPI Messages:	1.040e+02	1.00000	1.040e+02	2.080e+02
MPI Msg Lengths:	4.772e+05	1.00000	4.588e+03	9.544e+05
MPI Reductions:	1.450e+02	1.00000		



# Log summary: details

	Max	Ratio	Max	Ratio	Max	Ratio	Avg len	%T	%F	%M	%L	%R	%T	%F	%M	%L	%R	Mflop/s
MatMult	100	1.0	3.4934e-02	1.0	1.28e+08	1.0	8.0e+02	6	32	96	17	0	6	32	96	17	0	255
MatSolve	101	1.0	2.9381e-02	1.0	1.53e+08	1.0	0.0e+00	5	33	0	0	0	5	33	0	0	0	305
MatLUFactorNum	1	1.0	2.0621e-03	1.0	2.18e+07	1.0	0.0e+00	0	0	0	0	0	0	0	0	0	0	43
MatAssemblyBegin	1	1.0	2.8350e-03	1.1	0.00e+00	0.0	1.3e+05	0	0	3	83	1	0	0	3	83	1	0
MatAssemblyEnd	1	1.0	8.8258e-03	1.0	0.00e+00	0.0	4.0e+02	2	0	1	0	3	2	0	1	0	3	0
VecDot	101	1.0	8.3244e-03	1.2	1.43e+08	1.2	0.0e+00	1	7	0	0	35	1	7	0	0	35	243
KSPSetup	2	1.0	1.9123e-02	1.0	0.00e+00	0.0	0.0e+00	3	0	0	0	2	3	0	0	0	2	0
KSPSolve	1	1.0	1.4158e-01	1.0	9.70e+07	1.0	8.0e+02	26100	96	17	92	26100	96	17	92	26100	96	194

# User events

```
#include "petsclog.h"
int USER_EVENT;
PetscLogEventRegister(&USER_EVENT, "User event name", 0);
PetscLogEventBegin(USER_EVENT, 0, 0, 0, 0);
/* application code segment to monitor */
PetscLogFlops(number of flops for this code segment);
PetscLogEventEnd(USER_EVENT, 0, 0, 0, 0);
```